

Lyrical Analysis

2025-03-15

Load Libraries

```
library(class)
library(tm)
```

```
## Warning: package 'tm' was built under R version 4.3.3
## Loading required package: NLP
## Warning: package 'NLP' was built under R version 4.3.3
```

```
library(MASS)
library(nnet)
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.5.0      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
```

```
## -- Conflicts ----- tidyverse_conflicts() --
## x ggplot2::annotate() masks NLP::annotate()
## x dplyr::filter()      masks stats::filter()
## x dplyr::lag()          masks stats::lag()
## x dplyr::select()      masks MASS::select()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(tidyr)
library(caret)
```

```
## Warning: package 'caret' was built under R version 4.3.3
## Loading required package: lattice
##
## Attaching package: 'caret'
##
## The following object is masked from 'package:purrr':
##
##   lift
```

```
library(tidytext)
```

Load Data

```
# Read Lines
lines <- readLines("lyran.csv", encoding = "UTF-8", warn = FALSE)
## Incomplete Final Line Warning
# Get only valid lines
cleaned_lines <- iconv(lines, from = "UTF-8", to = "UTF-8", sub = "byte")
valid_lines <- cleaned_lines[!is.na(cleaned_lines)]
# Write Cleaned to New File
writeLines(valid_lines, "lyran_cleaned.csv")

# Read the Cleaned File
lyrics_data <- read.csv("lyran_cleaned.csv", fileEncoding = "UTF-8", stringsAsFactors = FALSE)
```

Stop Words and Cleaning Data

```
# Define your custom stop words
custom_stopwords <- c("ah", "im", "ohohoh", "oo", "uhhuh", "nooh", "s", "yearold", "thatll", "hadnt", "
# Combine custom stopwords with the default (english) stop words
all_stopwords <- c(stopwords("en"), custom_stopwords)
corpus <- Corpus(VectorSource(lyrics_data$lyrics))
# Clean the corpus
corpus_clean <- corpus %>%
  tm_map(tolower) %>%
  tm_map(removePunctuation) %>%
  tm_map(removeNumbers) %>%
  tm_map(removeWords, all_stopwords) %>%
  tm_map(stripWhitespace)

## Warning in tm_map.SimpleCorpus(., tolower): transformation drops documents
## Warning in tm_map.SimpleCorpus(., removePunctuation): transformation drops
## documents

## Warning in tm_map.SimpleCorpus(., removeNumbers): transformation drops
## documents

## Warning in tm_map.SimpleCorpus(., removeWords, all_stopwords): transformation
## drops documents

## Warning in tm_map.SimpleCorpus(., stripWhitespace): transformation drops
## documents

tidy_lyrics <- data.frame(
  name = rep(lyrics_data$name, each = sapply(corpus_clean, length)),
  artist = rep(lyrics_data$artist, each = sapply(corpus_clean, length)),
  theme = rep(lyrics_data$theme, each = sapply(corpus_clean, length)),
  text = sapply(corpus_clean, as.character),
  row_id = seq_along(lyrics_data$name), # Track the original row order
  stringsAsFactors = FALSE
)

## Warning in rep(lyrics_data$name, each = sapply(corpus_clean, length)): first
## element used of 'each' argument

## Warning in rep(lyrics_data$artist, each = sapply(corpus_clean, length)): first
```

```
## element used of 'each' argument

## Warning in rep(lyrics_data$theme, each = sapply(corpus_clean, length)): first
## element used of 'each' argument
```

Reformatting Data

```
# Make Unique Song_Artist_Id
lyrics_data <- tidy_lyrics %>%
  mutate(song_artist_id = paste(name, artist, sep = "_")) %>%
  dplyr::select(song_artist_id, theme, text)
```

Count Words

```
lyrics_data_tokens <- lyrics_data %>%
  mutate(text = as.character(text)) %>% # Ensure text is character
  unnest_tokens(output = word, input = text) %>% # Tokenize
  group_by(song_artist_id, theme, word) %>%
  summarise(word_count = n(), .groups = "drop") # Count words
```

Pivot

```
lyrics_data_wide <- lyrics_data_tokens %>%
  pivot_wider(names_from = word,
              values_from = word_count, # Use word_indicator instead of word_count
              values_fill = 0, # Fill missing values with 0
              names_glue = "word_{word}", # Custom column naming for words
              names_repair = "unique") # Make column names unique
```

Normalize Within Column

```
# Normalize each word column within the column (Min-Max Scaling)
lyrics_data_normalized <- lyrics_data_wide %>%
  mutate(across(starts_with("word_"), ~ (. - min(.)) / (max(.) - min(.)), .names = "norm_{.col}")) %>%
  dplyr::select(-starts_with("word_"))
```

Log Transform Then Do Within Column

```
lyrics_data_log_norm <- lyrics_data_wide %>%
  mutate(across(starts_with("word_"), ~ log(. + 1))) %>% # Log-transform
  mutate(across(starts_with("word_"), ~ (. - min(.)) / (max(.) - min(.)))) # Normalize
```

Pivot Theme

```
lyrics_theme_binary <- lyrics_data_tokens %>%
  group_by(song_artist_id, theme) %>%
  summarise(theme_indicator = 1, .groups = "drop")
```

```

lyrics_theme_onehot <- lyrics_theme_binary %>%
  pivot_wider(
    names_from = theme,
    values_from = theme_indicator,
    values_fill = 0,
    names_glue = "{theme}_theme" # Custom column naming for themes
  )

lyrics_data_log_norm <- as.data.frame(lyrics_data_log_norm)

lyrics_data_log_norm_factor <- lyrics_data_log_norm %>% mutate(theme_as_factor = factor(theme, ordered = TRUE))

lyrics_combined <- full_join(lyrics_theme_onehot, lyrics_data_log_norm_factor, by = join_by(song_artist_id, theme_as_factor))

```

Filter to Help Pick Words

```

lyrics_data_log_norm %>%
  filter(theme == "age/power dynamic") %>% # Filter for the 'age/power dynamic' theme
  pivot_longer(cols = -c(song_artist_id, theme), names_to = "word", values_to = "value") %>% # Pivot to long format
  filter(value > 0) %>% # Only include rows where the word appears (value > 0)
  group_by(song_artist_id, word) %>%
  summarise(word_value_in_song = sum(value), .groups = "drop") %>% # Count how many times the word appears in each song
  group_by(word) %>%
  summarise(
    total_word_value = sum(word_value_in_song), # Total word occurrences across all songs
    sd = sd(word_value_in_song),
    unique_songs = n_distinct(song_artist_id), # Number of unique songs the word appears in
    .groups = "drop"
  ) %>%
  filter(unique_songs > 1) %>%
  arrange(desc(total_word_value)) # Sort by total word count in descending order

```

```

## # A tibble: 358 x 4
##   word      total_word_value    sd unique_songs
##   <chr>          <dbl> <dbl>         <int>
## 1 word_never      8.29 0.188           19
## 2 word_know       7.99 0.196           18
## 3 word_now        6.70 0.165           17
## 4 word_just       6.66 0.137           16
## 5 word_like       5.65 0.153           18
## 6 word_love       5.52 0.184           13
## 7 word_baby       4.27 0.209           10
## 8 word_tell       4.06 0.167            9
## 9 word_cause      3.79 0.151           11
## 10 word_time      3.71 0.0898            8
## # i 348 more rows

```

Testing and Training Data

```

lyrics_data_log_norm_factor <- as.data.frame(lyrics_data_log_norm_factor)
lc_size <- floor(0.75 * nrow(lyrics_combined))

```

```
# Set Seed
set.seed(200)
train_ind_lc <- sample(seq_len(nrow(lyrics_combined)), size = lc_size)

train_lc <- lyrics_combined[train_ind_lc, ]
test_lc <- lyrics_combined[-train_ind_lc, ]

train_ldn <- lyrics_data_log_norm_factor[train_ind_lc, ]
test_ldn <- lyrics_data_log_norm_factor[-train_ind_lc, ]
```

LDA Fit

```
lda.fit <- lda(theme_as_factor ~ word_afford, data = lyrics_data_log_norm_factor, subset = train_ind_lc)
#print(lda.fit)
```

LDA Results

```
lda.pred <- predict(lda.fit, newdata = test_ldn)
lda.class <- lda.pred$class
lda_matrix <- table(Predicted = lda.class, Actual = test_ldn$theme_as_factor)
#print(lda_matrix)
actual_classes <- test_ldn$theme_as_factor
accuracy <- mean(lda.class == actual_classes)
print(paste("Accuracy:", round(accuracy * 100, 2), "%"))
```

```
## [1] "Accuracy: 3.18 %"
```

```
conf_matrix <- table(Predicted = lda.class, Actual = actual_classes)
#print(conf_matrix)
theme_accuracies <- diag(conf_matrix) / rowSums(conf_matrix)
theme_accuracies <- round(theme_accuracies * 100, 2) # Convert to percentage
print(theme_accuracies)
```

```
## age/power dynamic      crush      empowerment      exes
##           0.00          NaN          NaN          NaN
## forbidden love      grief      growing up      growth
##           NaN          NaN          NaN          NaN
## happy              hate      heartbreak      jealousy
##           NaN          NaN          NaN          NaN
## love      mental health      moving on      partying
##           NaN          NaN          3.23          NaN
## rebellion      religion      reminiscing      revenge
##           NaN          NaN          NaN          NaN
## situationship toxic relationship      unrequited
##           NaN          NaN          NaN
```

Testing Words

```
word_lm <- lm(word_yeah ~ `age/power dynamic_theme` + rebellion_theme + love_theme + `moving on_theme` + ...
anova_result <- anova(word_lm)
anova_summary <- broom::tidy(anova_result)
```

```
# Filter only significant variables (p-value < 0.05)
significant_results <- anova_summary %>%
  arrange(p.value) %>%
  filter(p.value < 0.0001)
# Display results
print(significant_results)
```

```
## # A tibble: 2 x 6
##   term                df sumsq meansq statistic    p.value
##   <chr>              <int> <dbl>  <dbl>    <dbl>    <dbl>
## 1 happy_theme        1 0.704  0.704    17.2 0.0000381
## 2 partying_theme     1 0.628  0.628    15.4 0.0000986
```

Feature Selection

Correlation Based Feature Selection

```
#X <- train_lc[, -c(1:24, 7332)]
#Y <- train_lc$theme_as_factor
```

Recursive Feature Elimination

```
#control <- rfeControl(functions = rfFuncs, method = "cv", number = 5)
#rfe_result <- rfe(X, Y, sizes = c(1:10), rfeControl = control)

# Selected features
#X_selected <- X[, predictors(rfe_result)]

# Remove them
#X_filtered <- X[, -highly_correlated]
```

Multi Log

```
multi_log_model <- multinom(theme_as_factor ~ word_love + word_baby +
  word_yeah + word_still + word_never,
  data = train_lc, trace = FALSE)
multi_log_pred <- predict(multi_log_model)#, newdata = test_lc)
# Calculate Accuracy
accuracy <- mean(multi_log_pred == train_lc$theme_as_factor)
print(paste("Multinomial Logistic Regression Accuracy:", round(accuracy * 100, 2), "%"))
```

```
## [1] "Multinomial Logistic Regression Accuracy: 14.71 %"

# Create Confusion Matrix
conf_matrix <- table(Predicted = multi_log_pred, Actual = train_lc$theme_as_factor)
# Compute Accuracy per Theme (Diagonal / Row Sum)
theme_accuracies <- diag(conf_matrix) / rowSums(conf_matrix)
theme_accuracies <- round(theme_accuracies * 100, 2) # Convert to percentages
# Print Theme-Wise Accuracy
print(theme_accuracies)
```

```
## age/power dynamic      crush      empowerment      exes
```

```
##          33.33          19.48          31.58          0.00
## forbidden love      grief      growing up      growth
##          0.00          NaN          6.67          0.00
##          happy      hate      heartbreak      jealousy
##          0.00          0.00          23.08          0.00
##          love      mental health      moving on      partying
##          16.25          7.69          9.88          11.11
##          rebellion      religion      reminiscing      revenge
##          8.70          NaN          26.67          0.00
##          situationship toxic relationship      unrequited
##          20.00          0.00          13.64
```

```
# Calculate column sums
```

```
word_sums <- colSums(train_lc[, grepl("^word_", colnames(train_lc))])
```

```
# Select columns where sum is at least 2
```

```
selected_words <- names(word_sums[word_sums >= 35])
```

```
# Subset dataset with only selected word features + target variable
```

```
train_lc_filtered <- train_lc[, c("theme_as_factor", selected_words)]
```

```
#full_model <- multinom(theme_as_factor ~ ., data = train_lc_filtered, trace = FALSE)
```

```
#step_model <- stepAIC(full_model, direction = "backward", trace = TRUE)
```

```
#summary(step_model) # View selected variables
```